

Machine Learning



Machine Learning

Lecture: Instance-Based Learning

Ted Scully

Instance Based Learning

- ▶ Instance-based learning is a family of learning algorithms that **compare new problem instances with existing instances in the training data.**
- ▶ Predictions for new instances are based on their **similarity to stored instances** (the basis of the similarity measure is typically distance)

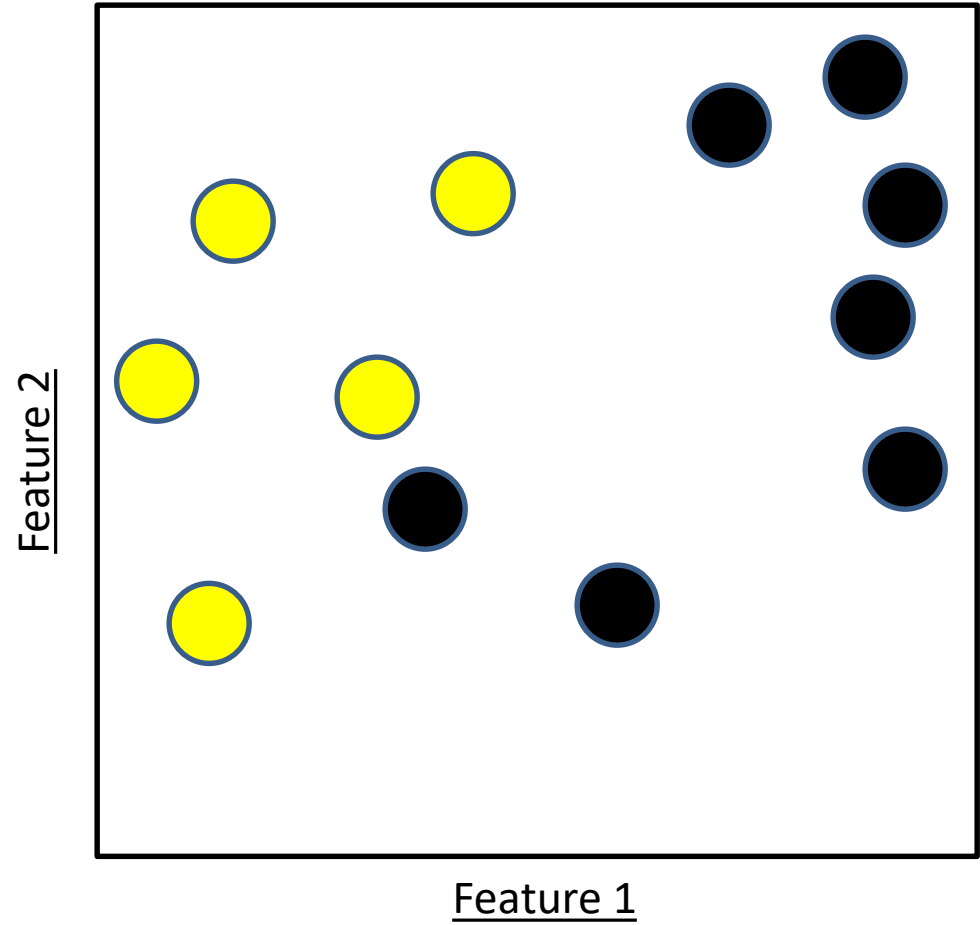
Nearest Neighbour Algorithm (1)

- ▶ The Nearest Neighbour algorithm is the simplest form of IBL
- ▶ Nearest Neighbour algorithm:
 - ▶ Given a test case with a value to be predicted, identify which stored case it is nearest.
 - ▶ Assigns the **new test case the same class as the nearest neighbour**
 - ▶ Requires a distance metric.
- ▶ This very simple algorithm is very susceptible to noise.

Given a query instance \mathbf{x}_q ,
first locate the nearest training example \mathbf{x}_n
then $\mathbf{f}(\mathbf{x}_q) := \mathbf{f}(\mathbf{x}_n)$


Where $\mathbf{f}(\mathbf{x}_n)$ is the class associated with the data item \mathbf{x}_n

| Feature 1 | Feature 2 | Colour |
|-----------|-----------|--------|
| | | |
| | | |
| | | |
| | | |
| | | |
| | | |
| | | |
| | | |
| | | |
| | | |
| | | |



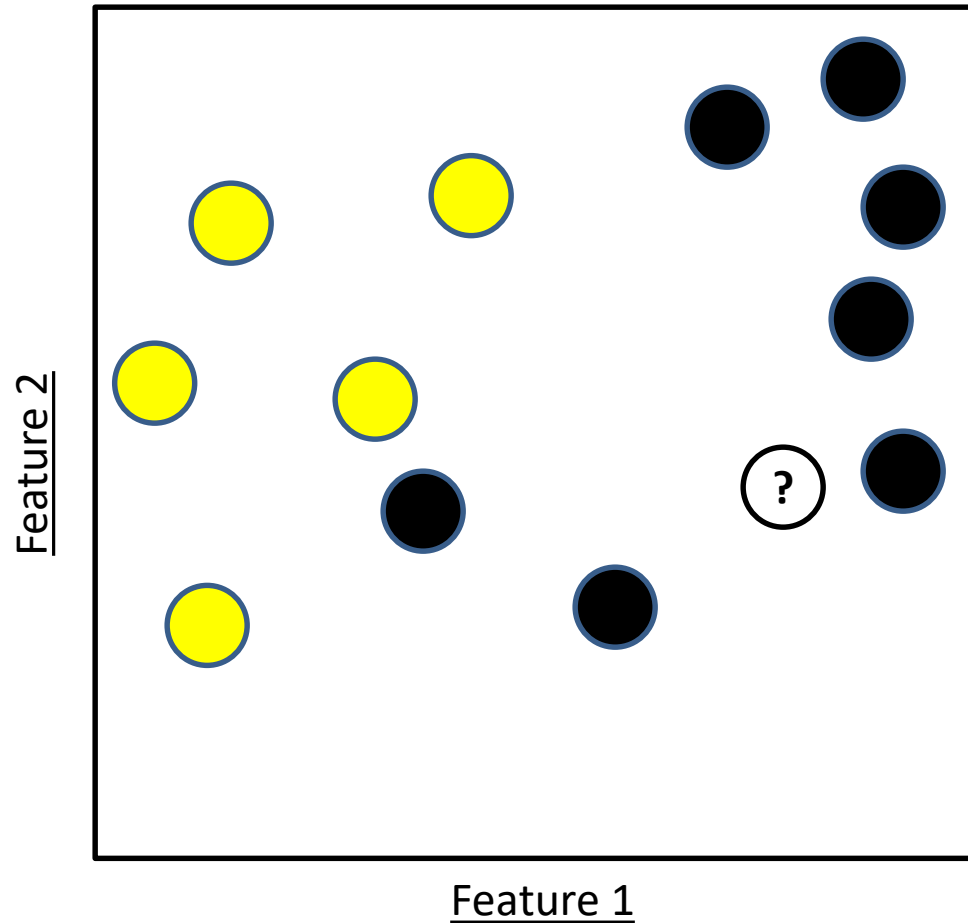
| Feature 1 | Feature 2 | Colour |
|-----------|-----------|--------|
| | | |
| | | |
| | | |
| | | |
| | | |
| | | |
| | | |
| | | |

- ▶ We can represent a dataset in an IBL by mapping all instances to a **feature space**, that is, using each descriptive feature as an axis of a coordinate system.
- ▶ We can then place each instance within the feature space based on the value of it's features.




Nearest Neighbour Example

- ▶ We wish to classify the new case, which is the white circle with the question mark.
- ▶ The nearest neighbour in feature space is selected and the example instance is assigned the same class.



| Feature 1 | Feature 2 | Regression Target |
|-----------|-----------|-------------------|
| | | |
| | | |
| | | |
| | | |
| | | |
| | | |
| | | |
| | | |
| | | |
| | | |
| | | |

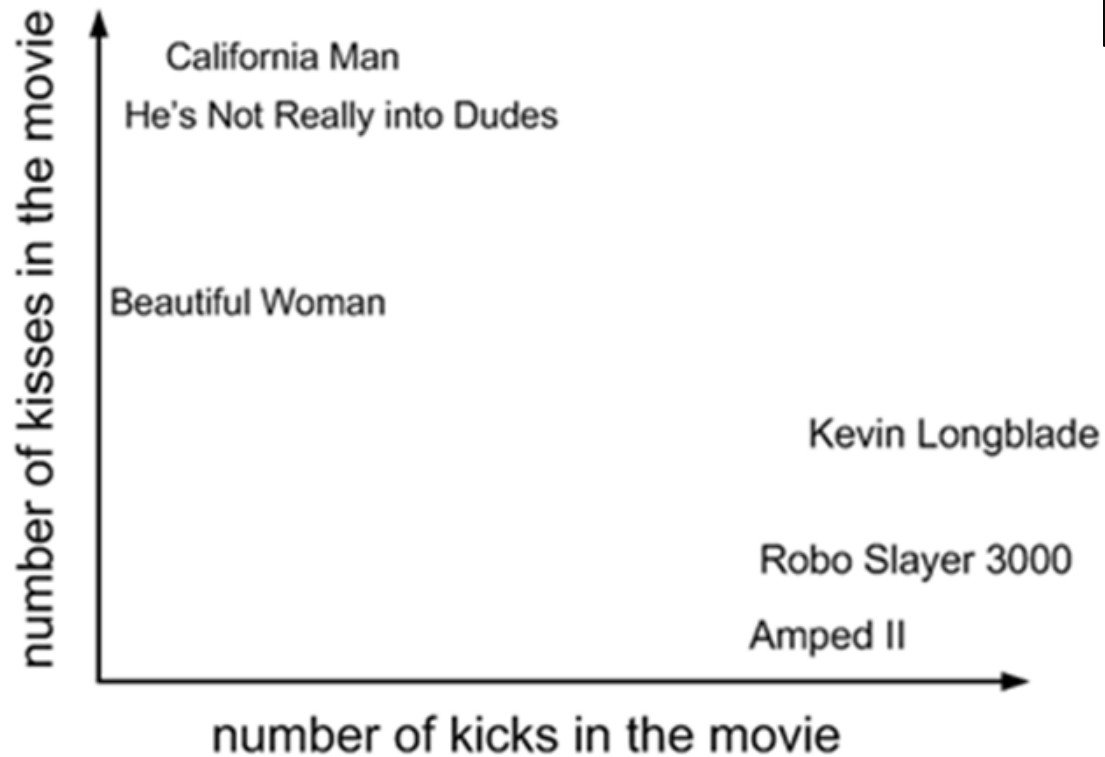


Sample Dataset

| Movie title | # of kicks | # of kisses | Type of movie |
|-----------------------------------|------------|-------------|---------------|
| <i>California Man</i> | 3 | 104 | Romance |
| <i>He's Not Really into Dudes</i> | 2 | 100 | Romance |
| <i>Beautiful Woman</i> | 1 | 81 | Romance |
| <i>Kevin Longblade</i> | 101 | 10 | Action |
| <i>Robo Slayer 3000</i> | 99 | 5 | Action |
| <i>Amped II</i> | 98 | 2 | Action |

Feature Space

Here we can see the feature space for our simple movie dataset.

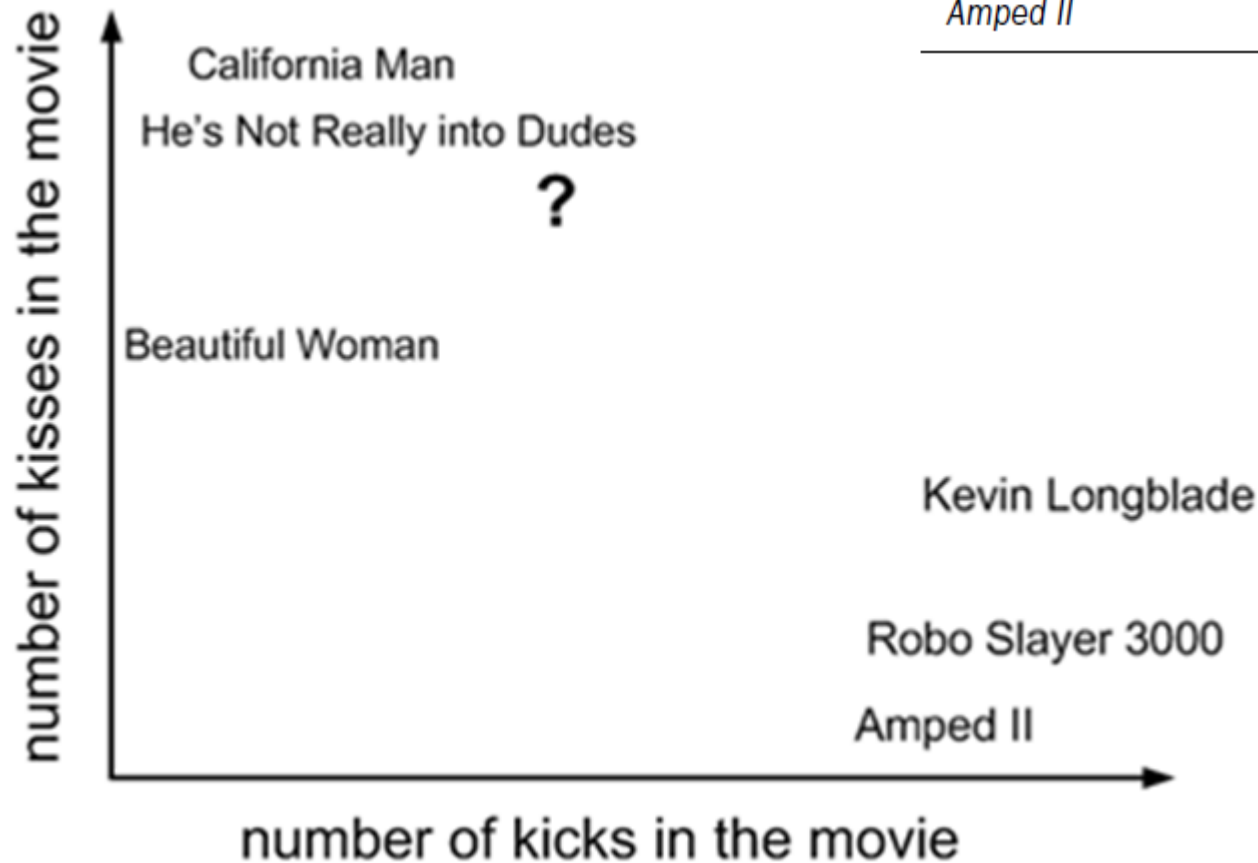


| Movie title | # of klicks | # of kisses | Type of movie |
|-----------------------------------|-------------|-------------|---------------|
| <i>California Man</i> | 3 | 104 | Romance |
| <i>He's Not Really into Dudes</i> | 2 | 100 | Romance |
| <i>Beautiful Woman</i> | 1 | 81 | Romance |
| <i>Kevin Longblade</i> | 101 | 10 | Action |
| <i>Robo Slayer 3000</i> | 99 | 5 | Action |
| <i>Amped II</i> | 98 | 2 | Action |
| ? | 18 | 90 | Unknown |

Assume we get an unseen movie and we have to classify it as a Romance or action based on it's feature values.

Given a query instance \mathbf{x}_q ,
 first locate the nearest
 training example \mathbf{x}_n
 then $\mathbf{f}(\mathbf{x}_q) := \mathbf{f}(\mathbf{x}_n)$

| Movie title | Distance to movie “?” |
|----------------------------|-----------------------|
| California Man | 20.5 |
| He’s Not Really into Dudes | 18.7 |
| Beautiful Woman | 19.2 |
| Kevin Longblade | 115.3 |
| Robo Slayer 3000 | 117.4 |
| Amped II | 118.9 |



As the query is
 closest to the
 film “He’s Not
 Really into
 Dudes” then it is
 classified as a
Romance

| Movie title | # of kicks | # of kisses | Type of movie |
|-----------------------------------|------------|-------------|---------------|
| <i>California Man</i> | 3 | 104 | Romance |
| <i>He's Not Really into Dudes</i> | 2 | 100 | Action |
| <i>Beautiful Woman</i> | 1 | 81 | Romance |
| <i>Kevin Longblade</i> | 101 | 10 | Action |
| <i>Robo Slayer 3000</i> | 99 | 5 | Action |
| <i>Amped II</i> | 98 | 2 | Action |
| ? | 18 | 90 | Unknown |

Noise: An incorrect classification



What would happen if “He’s no really into dudes” was incorrectly classified as an Action movie.

Our new query instance ‘?’ would also get **incorrectly classified** as an Action.

The simple nearest neighbour approach is very likely to over-fit on the training data.

Any ideas of how we might go about improving the algorithm?

K-Nearest Neighbour


- ▶ A simple extension is to consider not just the nearest neighbour, but **several nearest neighbours**.
- ▶ This requires defining a neighbourhood; the standard approach is to use a neighbourhood that is just large enough to include a fixed number of points, k .
- ▶ But how do we decide on **appropriate class** or **regression value** if we are consider more than one neighbour in feature space???

K-Nearest Neighbour

- ▶ A simple extension is to consider not just the nearest neighbour, but **several nearest neighbours**.
- ▶ This requires defining a neighbourhood; the standard approach is to use a neighbourhood that is just large enough to include a fixed number of points, k .
- ▶ Prediction is based on these k nearest neighbours.
 - ▶ If this is a **regression** problem then use the **average** of k -nearest neighbours.
 - ▶ If it is a **classification** problem then take a **vote** amongst the k -nearest neighbours
 - ▶ This approach is less sensitive to noise.

k - Nearest Neighbour Algorithm - Classification

- ▶ k-NN can be applied to classification problems
- ▶ If it is a classification problem then take a **majority vote** amongst the k-nearest neighbours
- ▶ What is the classification of the query instance if $k = 3$? What if $k = 5$



k - Nearest Neighbour Algorithm - Regression

We assume a set of training examples $\langle x_i, f(x_i) \rangle$

Given a query instance x_q ,


Identify k nearest training examples

If it is regression problem, then average values of k -nearest neighbours

$$f(x_q) := \frac{\sum_{i=1}^k f(x_i)}{k}$$



k - Nearest Neighbour Algorithm - Regression

- ▶ k – NN can be applied to regression problems
- ▶ The answer is the average value of k of the neighbours
- ▶ What is the answer if: $k = 3$?
 $(82+78+83)/3 = 81.$



Selecting an Appropriate K Value


- ▶ The selection of an appropriate value of k is very important. The parameter k is what we refer to as a hyper-parameter
- ▶ Selecting **too small** a value can make the algorithm **susceptible to noise** and can **overfit on the training data**.
- ▶ Selecting larger values of k lessen the impact of noise on the classification but make **boundaries between classes less distinct (in other words the model can underfit and the boundary fails to capture the patterns in the data)**.
- ▶ There are many techniques for selecting a k value.
 - ▶ Certain rules-of-thumb such as use the square root of the number of classified instances [not recommended].
 - ▶ Instead you should select a **range of different k values** and assess the performance of your model for these values (later when we cover scikit learn we will look at using N-fold cross validation and search to identify good values for k)



Distance Metrics

- ▶ An important aspect of k-NN algorithms is how we determine which instances are the nearest to the target case. Thus, the distance metric is a measure of the similarity between two cases.
- ▶ Common distance metrics include:
 - ▶ Euclidean
 - ▶ Manhattan
 - ▶ Minkowski

Distance Metrics - Euclidean



- ▶ To help illustrate the various metrics let's assume we have the dataset below with n features and two instances p and q

| | Feature 1 | Features 2 | | Feature n |
|-----|-----------|------------|-------|-------------|
| p | p_1 | p_2 | | p_n |
| q | q_1 | q_2 | | q_n |

Euclidean Distance Metric

- ▶ If $\mathbf{p} = \langle p_1, p_2, \dots, p_n \rangle$ and $\mathbf{q} = \langle q_1, q_2, \dots, q_n \rangle$ are two points in Euclidean n -space, then the distance (d) from \mathbf{p} to \mathbf{q} , or from \mathbf{q} to \mathbf{p} is given by

$$\begin{aligned} d(\mathbf{p}, \mathbf{q}) &= d(\mathbf{q}, \mathbf{p}) = \sqrt{(q_1 - p_1)^2 + (q_2 - p_2)^2 + \dots + (q_n - p_n)^2} \\ &= \sqrt{\sum_{i=1}^n (q_i - p_i)^2}. \end{aligned}$$

It is important to understand that \mathbf{p} and \mathbf{q} here represent two data instances. Each instance consisting of a finite set of features. The instance \mathbf{q} has the features q_1, q_2, \dots, q_n .

Euclidean Distance Metric

| Movie title | # of kicks | # of kisses | Movie title | Distance to movie “?” |
|----------------------------|------------|-------------|----------------------------|-----------------------|
| California Man | 3 | 104 | California Man | 20.5 |
| He's Not Really into Dudes | 2 | 100 | He's Not Really into Dudes | 18.7 |
| Beautiful Woman | 1 | 81 | Beautiful Woman | 19.2 |
| Kevin Longblade | 101 | 10 | Kevin Longblade | 115.3 |
| Robo Slayer 3000 | 99 | 5 | Robo Slayer 3000 | 117.4 |
| Amped II | 98 | 2 | Amped II | 118.9 |
| ? | 18 | 90 | Action | |
| | | | Unknown | |

Distance between **California man** (3, 104) and the **query** instance (18, 90) would be:

Euclidean Distance Metric


| Movie title | # of clicks | # of classes | Movie title | Distance to movie “?” |
|----------------------------|-------------|--------------|----------------------------|-----------------------|
| California Man | 3 | 104 | California Man | 20.5 |
| He’s Not Really into Dudes | 2 | 100 | He’s Not Really into Dudes | 18.7 |
| Beautiful Woman | 1 | 81 | Beautiful Woman | 19.2 |
| Kevin Longblade | 101 | 10 | Kevin Longblade | 115.3 |
| Robo Slayer 3000 | 99 | 5 | Robo Slayer 3000 | 117.4 |
| Amped II | 98 | 2 | Amped II | 118.9 |
| ? | 18 | 90 | Action | |
| | | | Unknown | |

Distance between **California man** (3, 104) and the **query** instance (18, 90) would be

$$\sqrt{(3 - 18)^2 + (104 - 90)^2} = \sqrt{225+196}= 20.5$$


Manhattan Distance Metric

- ▶ Manhattan distance measures distance **parallel to each axis**, not diagonally (in downtown Manhattan, to get from one point to another you generally walk North-South and East-West, rather than 'as the crow flies').
- ▶ In other words take the sum of the absolute values of the differences of the coordinates
- ▶ $d(p, q) = |q_1 - p_1| + |q_2 - p_2| + |q_n - p_n| = \sum_{i=1}^n |q_i - p_i|$




Manhattan Distance Metric

- ▶ Lets assume we have a simple dataset containing three instances as follows. We are also given the query instance below. Calculate the Manhattan and Euclidean distance between the **first training example** and the **query**



Manhattan



Euclidean

| Area | Weight | Height | Capacity |
|------|--------|--------|----------|
| 10 | 8 | 4 | 14 |
| 12 | 10 | 6 | 12 |
| 14 | 9 | 4 | 11 |

| Area | Weight | Height | Capacity |
|------|--------|--------|----------|
| 5 | 4 | 4 | 10 |

| Area | Weight | Height | Capacity |
|------|--------|--------|----------|
| 10 | 8 | 4 | 8 |
| 12 | 10 | 6 | 12 |
| 14 | 9 | 4 | 11 |

| Area | Weight | Height | Capacity |
|------|--------|--------|----------|
| 5 | 4 | 4 | 3 |

Manhattan Distance Metric

- ▶ Euclidean

- ▶ $(10-5)^2 + (8-4)^2 + (4-4)^2 + (8-3)^2$

- ▶ 57

- ▶ Square root of 57 = 7.55

- ▶ Manhattan

- ▶ $|10-5| + |8-4| + |4-4| + |8-3|$

- ▶ $5 + 4 + 0 + 5 = 14$

| Area | Weight | Height | Capacity |
|------|--------|--------|----------|
| 10 | 8 | 4 | 8 |
| 12 | 10 | 6 | 12 |
| 14 | 9 | 4 | 11 |

| Area | Weight | Height | Capacity |
|------|--------|--------|----------|
| 5 | 4 | 4 | 3 |

Minkowski Distance

- ▶ The Minkowski distance between a feature vector $\mathbf{p} = \langle p_1, p_2, \dots, p_n \rangle$ and another feature vector $\mathbf{q} = \langle q_1, q_2, \dots, q_n \rangle$ is defined as:

- ▶
$$d(p, q) = (\sum_{i=1}^n |p_i - q_i|^a)^{\frac{1}{a}}$$

- ▶ In the above equation a is an integer. Consider the case when $a = 1$ or $a = 2$. Any comments.

Minkowski Distance

- ▶ The Minkowski distance between a feature vector $\mathbf{p} = \langle p_1, p_2, \dots, p_n \rangle$ and another feature vector $\mathbf{q} = \langle q_1, q_2, \dots, q_n \rangle$ is defined as:
- ▶
$$d(p, q) = (\sum_{i=1}^n |p_i - q_i|^a)^{\frac{1}{a}}$$
- ▶ In the above equation a is an integer. Consider the case when $a = 1$ or $a = 2$. Any comments.
 - ▶ For $a = 1$ we get the Manhattan distance and for $a = 2$ we get the Euclidean distance.
 - ▶ Minkowski Distance is a generalization of the Euclidean and Manhattan distance metrics.

$$d(p, q) = \left(\sum_{i=1}^n |p_i - q_i|^a \right)^{\frac{1}{a}}$$

Larger values of a place more emphasis on large differences between feature values compare to smaller values. This is because the differences are raised to the power of a . Therefore, the Euclidean distance weights features with larger differences between feature value influence the final distance metric more than features with a smaller difference.

$$d(p, q) = \left(\sum_{i=1}^n |p_i - q_i|^a \right)^{\frac{1}{a}}$$

As you might expect as you begin to decrease a the opposite effect happens. The features with larger differences don't have the same level of impact.

Building a K Nearest Neighbour Classifier

- ▶ **Step 1. Read information from a dataset**
 - ▶ Read data from a dataset containing classified instances. Read each feature of the dataset as well as corresponding class.
- ▶ **Step 2. Determine distance between each dataset entry and the query instance**
 - ▶ Use a suitable distance metric to calculate the distance between the query instance and all k neighbours
- ▶ **Step 3. Classify the query instance**
 - ▶ Identify k nearest data instances. Assign query instance category corresponding to most common category.

Problems Measuring Distance 1 -Scale

- ▶ Since the performance of k-NN is strongly dependent on the choice of distance metric, you need to be aware of some pitfalls.
- ▶ The first problem arises when the features are different from each other.
- ▶ For example, if one feature has a range between **0 and 1** and another feature has a range between **0 and 10, 000**, it hardly makes sense to add them as would happen with Euclidian or Manhattan distance metrics (for example, salary and age).
- ▶ What is the main problem that arises from the above situation?

Problems Measuring Distance (1)

- ▶ **Problem 1: Scaling**

- ▶ Feature A has range 1-10
Feature B has range 1-1000
- ▶ Feature B will dominate calculations

- ▶ Example, lets calculate the distance between data instance 1 and 2
 - ▶ Data instance 1 = (5.5, 787)
 - ▶ Data instance 2 = (7.5, 567)

Problems Measuring Distance (1)

- ▶ **Problem 1: Scaling**

- ▶ Feature A has range 1-10
Feature B has range 1-1000
- ▶ Feature B will dominate calculations

- ▶ Example, lets calculate the distance between data instance 1 and 2
 - ▶ Data instance 1 = (5.5, 787)
 - ▶ Data instance 2 = (7.5, 567)

$$\sqrt{(5.5 - 7.5)^2 + (787 - 567)^2}$$

$$\sqrt{16 + 48400}$$

Problems Measuring Distance (1)

▶ Problem 1: Scaling

- ▶ Feature A has range 1-10
Feature B has range 1-1000
- ▶ Feature B will dominate calculation

We can see below that the second feature is entirely dominating the distance calculation simply because it has a larger range of values compared to the first feature.

▶ Example, lets calculate the distance

- ▶ Data instance 1 = (5.5, 787)
- ▶ Data instance 2 = (7.5, 567)

We don't want our model to bias toward a particular feature simply because the range happens to be larger.

$$\sqrt{(5.5 - 7.5)^2 + (787 - 567)^2}$$

$$\sqrt{16 + 48400}$$

Problems Measuring Distance (1)

- ▶ Solution:
 - ▶ Normalise all dimensions independently (scale data so that it has a maximum and minimum range)
 - ▶ Using range normalization we identify the minimum and maximum value for a specific feature. We can then apply the following formula.
 - ▶
$$newValue = \frac{originalValue - minValue}{maxValue - minValue}$$

$$\triangleright \text{newValue} = \frac{\text{originalValue} - \text{minValue}}{\text{maxValue} - \text{minValue}}$$

- ▶ Problem 1: Scaling
 - ▶ Feature A has range 1-10
Feature B has range 1-1000
- ▶ Normalise variables
 - ▶ Feature A
 - ▶ Feature B

- ▶ Before Normalization
 - ▶ Data instance 1 = (5.5, 787)
 - ▶ Data instance 2 = (7.5, 567)

$$\text{newValue} = \frac{\text{originalValue} - \text{minValue}}{\text{maxValue} - \text{minValue}}$$

- ▶ Problem 1: Scaling

- ▶ Feature A has range 1-10
 - Feature B has range 1-1000

- ▶ Normalise variables

- ▶ Feature A

- ▶ $(5.5 - 1)/(10-1) = 0.5$
 - ▶ $(7.5 - 1)/(10 - 1) = 0.72$

- ▶ Feature B

- ▶ $(787-1)/(1000-1) = 0.78$
 - ▶ $(567-1)/(1000-1) = 0.56$

- ▶ Before Normalization

- ▶ Data instance 1 = (5.5, 787)
 - ▶ Data instance 2 = (7.5, 567)

- ▶ After Normalization

- ▶ Data instance 1 = (0.5, .78)
 - ▶ Data instance 2 = (0.72, 0.56)

$$\text{newValue} = \frac{\text{originalValue} - \text{minValue}}{\text{maxValue} - \text{minValue}}$$

▶ Problem 1: Scaling

- ▶ Feature A has range 1-10
Feature B has range 1-1000

▶ Normalise variables

- ▶ Feature A
 - ▶ $(5.5 - 1)/(10-1) = 0.5$
 - ▶ $(7.5 - 1)/(10 - 1) = 0.72$
- ▶ Feature B
 - ▶ $(787-1)/(1000-1) = 0.78$
 - ▶ $(567-1)/(1000-1) = 0.56$

▶ Before Normalization

- ▶ Data instance 1 = (5.5, 787)
- ▶ Data instance 2 = (7.5, 567)

▶ After Normalization

- ▶ Data instance 1 = (0.5, .78)
- ▶ Data instance 2 = (0.72, 0.56)

$$\sqrt{(0.5 - 0.72)^2 + (0.78 - 0.56)^2}$$

$$\sqrt{0.048 + 0.048}$$

Problems Measuring Distance (1)

- ▶ When we normalize the train data, it is also important to understand:
 - ▶ We normalize each feature **independently**
 - ▶ We must normalize the test data using the same parameters for max and min (that is we still use the minValue and maxValue from the original training set).

Problems Measuring Distance – Irrelevant Features

- ▶ The other principal problem is that **all features are included equally** in the calculations we have looked at, even though some features may be **redundant** or **less relevant**.
- ▶ Therefore, a number of features may skew the result even though they might have little or no impact on the classification.
- ▶ **Solution 2A:**
 - ▶ **Assign weighting** to each dimension (Optimise weighting to minimise error)
- ▶ **Solution 2B:**
 - ▶ Give some dimensions **0 weight** (Feature subset solution)
- ▶ Either way, since we cannot know in advance what weighting to give dimensions, systematic repeated experiments are needed to optimise them.
- ▶ We will look at feature selection in more detail later in the module.